

Rockborne

PySpark: Setup, Essentials, and Data Transformation

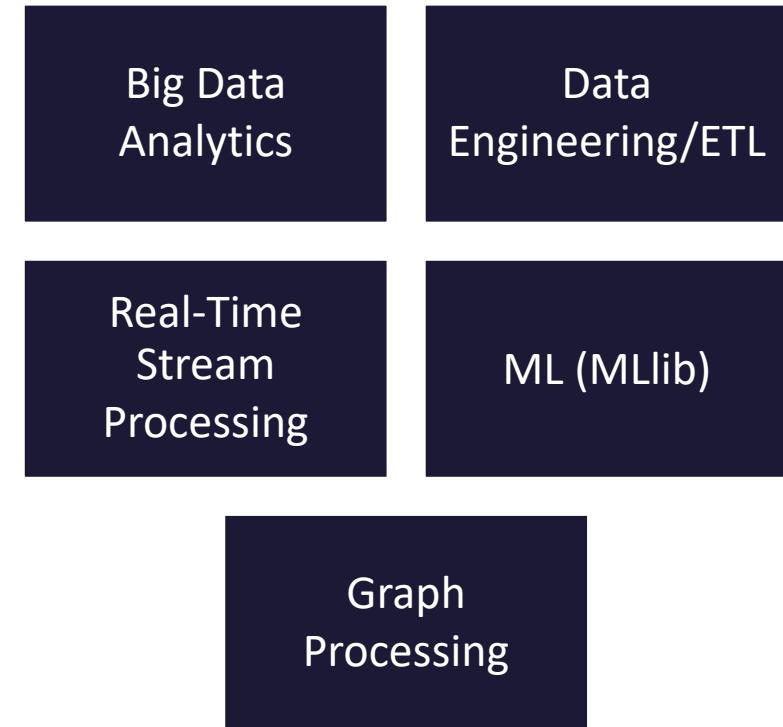
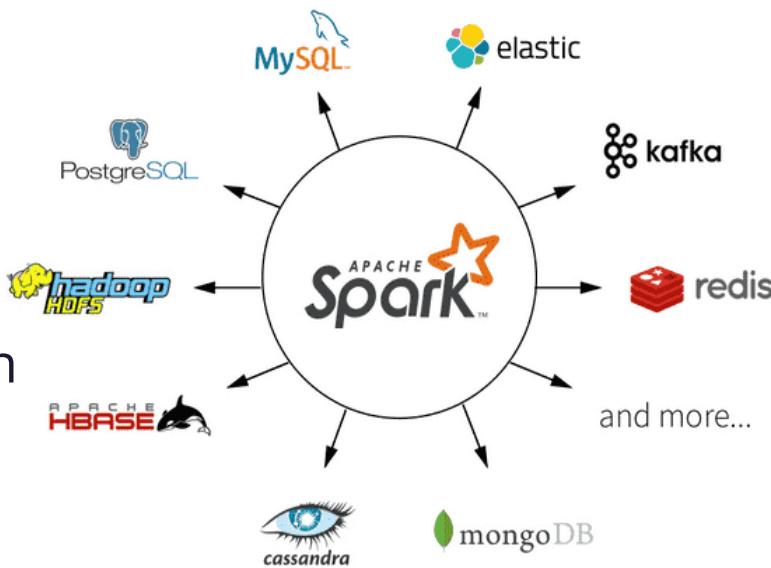
Workshop Outline

- PySpark Essentials
 - Introduction to PySpark
 - Setting up the environment
 - Basic operations
 - RDDs and Data frames
 - Data frames and SQL
- Data Transformation with PySpark
 - Filtering
 - Manipulation
 - Aggregation
 - JOINs and UNION



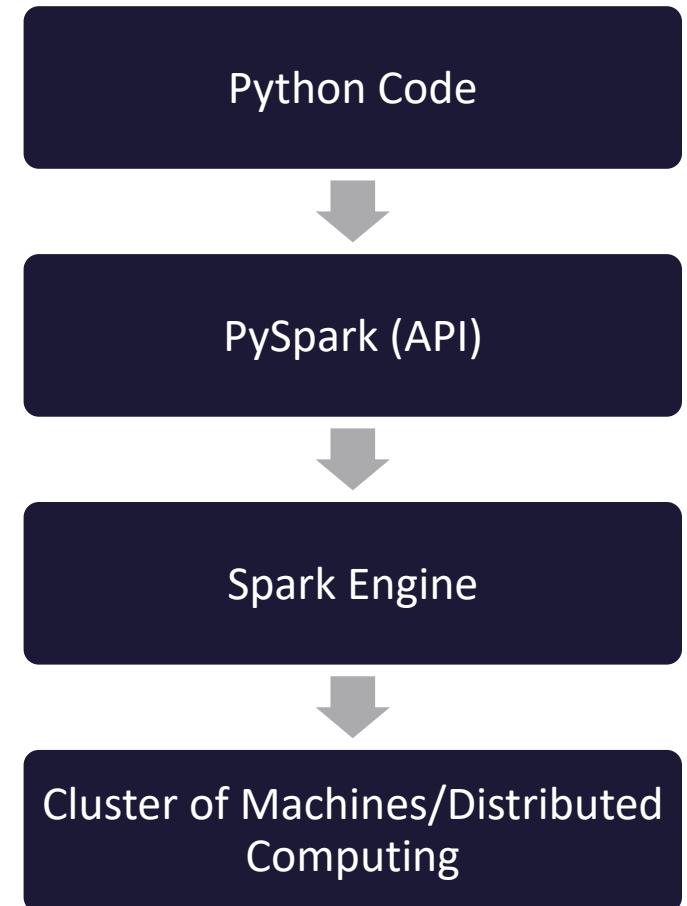
Apache Spark

- Apache Spark is an open source, distributed computing system
- It processes large-scale data across multiple machines (clusters)
- Provides in-memory computation
 - Tasks are completed in place in memory, increasing energy efficiency
- Supports multiple programming languages



PySpark

- PySpark is the Python API for Apache Spark
- It allows you to write Spark applications using Python while still provides access to Spark's core features:
 - RDD API: low-level distributed data analysis
 - DataFrame API: high-level, optimised SQL-like tables
 - SQL API: query data with SWL syntax
- It integrates easily with the Python ecosystem (NumPy, Pandas, etc.)



PySpark Essentials:

- Basic Operations
 - RDDs and DFs
 - DFs and SQL

Why PySpark?

PySpark is the Python API for Apache Spark, allowing for distributed data processing:

- It handles big datasets across multiple machines
 - It splits large datasets into smaller chunks for parallel processing
 - Scales to terabytes of data across clusters
- Provides both low-level control (RDDs) and high-level abstraction (Dataframes)
 - RDDs: fine-grained control, custom transformations
 - Dataframes: SQL-like and user-friendly for most analytics
- Integrates easily with Python
 - Works in Jupyter Notebooks
 - Converts well between Spark and Pandas Dataframes
 - Enables end-to-end data analysis and ML pipelines

Spark

Analogy:



Driver The head chef who plans the menu.



Executors The line cooks who prepare dishes.



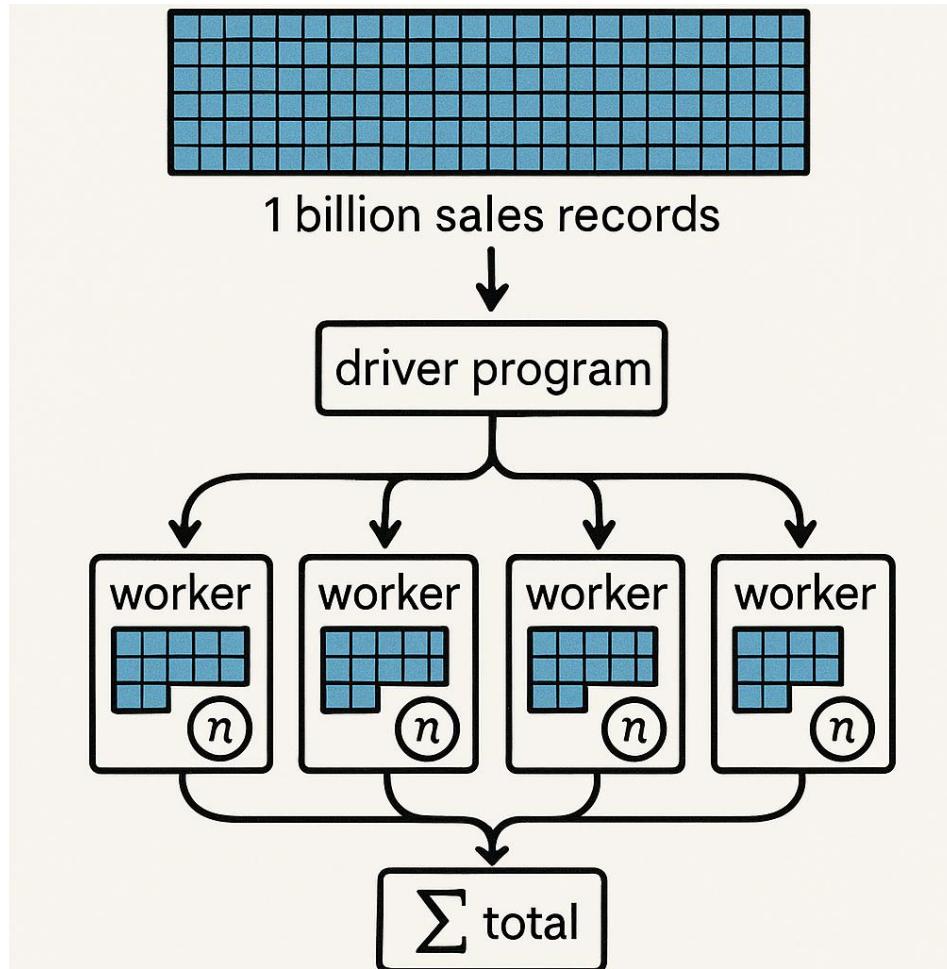
Cluster Manager The restaurant manager who assigns resources

sends tasks



Results come back

Spark



Basic Operations using PySpark

Loading and Inspecting data:

1. The first step is to always load your data into a Dataframe.

- o This allows for external datasets to be distributed and processed

```
df = spark.read.csv("superstore.csv", header=True, inferSchema=True)
```

2. Inspecting the data allows you to quickly explore and understand the data structure and quality

- o Ensures correct schema and makes catching data issues easy

```
df.show(5) ## Shows first 5 rows
```

```
df.printSchema() ## Print schema with column types
```

```
df.columns ##Column names
```

Basic Operations using PySpark

Dimensions, and Statistics

1. Unlike Pandas, Spark separates rows and columns

- This is important for large datasets

```
df.count() ##Row count
```

```
len(df.columns) ##Column count
```

- Spark also allows you to get a quick summary of numerical fields for initial exploration

```
df.describe().show() ##Summary statistics for all numerical columns
```

2. Spark allows you to narrow down to the columns you need, and focus on subsets of data

- This improves performance and readability
- Common in data cleaning and analysis

```
df.select("Sales").show(5) ##Single column
```

```
df.select("Category", "Sales").show(5) ##Multiple columns
```

```
df.filter(df["Category"] == "Furniture").show(5) ##Equality filter
```

```
df.filter(df["Sales"] > 500).show(5) ##Numeric condition
```

RDDs and DataFrames using PySpark

Resilient Distributed Datasets (RDDs) and Transformations

- They are low-level distributed collection of objects , and are the foundation of Spark
 - They allow for custom transformations and fine-grained control

```
rdd = df.rdd ##Convert df to RDD  
rdd.take(5) ##Take sample rows
```
- They transform data only when an action is called, allowing for parallel and optimised execution.
 - Without these actions being called, Spark won't compute anything

```
sales_rdd = rdd.map(lambda row: row.Sales)  
  
high_sales = sales_rdd.filter(lambda x: float(x) > 500)
```

```
sales_rdd.collect() ##Collect as list
```

```
sales_rdd.take(5) ##First 5
```

```
rdd.count() ##Count rows
```

RDDs and DataFrames using PySpark

RDD Actions

- Actions trigger execution and return results, and without these, Spark won't compute anything

```
sales_rdd.collect() ##Collect as list
```

RDDs to Dataframes:

- RDDs are flexible, but dataframes are easier to use

- They give SQL-like power and optimisations

```
from pyspark.sql import Row
```

```
df2 = rdd.map(lambda row: Row(Sales=row.Sales)).toDF() ##Convert RDD back to DF
```

```
df2.show(5)
```

```
sales_rdd.take(5) ##First 5
```

```
rdd.count() ##Count rows
```

RDDs and DataFrames using PySpark

Dataframes

- They are a distributed table with named columns (similar to Pandas DF)

- They are the most common Spark API, with easy syntax and optimised execution

```
df.show(5) ##df already created from csv
```

```
df.printSchema() ##Check schema
```

- You can add, rename, or transform columns – all of which are core steps

```
from pyspark.sql.functions import col
```

```
df = df.withColumn("Sales_x2", col("Sales") * 2) ##Add a new column
```

```
df = df.withColumnRenamed("Sales", "Total_Sales") ##Rename column
```

RDDs and DataFrames using PySpark

Grouping and Aggregation

- Grouping and aggregation allows you to summarise by categories, which is essential for analysis and business ins

```
from pyspark.sql.functions import col  
df = df.withColumn("Sales_x2", col("Sales") * 2) ##Add a new column  
df = df.withColumnRenamed("Sales", "Total_Sales") ##Rename column
```

Using SQL with Dataframes

- Registering the dataframe as a table allows analysts to reuse SQL skills in Spark

```
df.createOrReplaceTempView("orders") ## Register table  
  
spark.sql("""  
    SELECT Category, SUM(Total_Sales)  
    FROM orders  
    GROUP BY Category  
""").show() ## Run SQL query
```

RDDs and DataFrames using PySpark

Filtering

- You should always only keep rows that meet your project's conditions
 - This removes noise and allows you to focus on relevant data

```
From pyspark.sql.functions import col
```

```
Df.filter(col("Sales") > 500).show() ##Filter rows where sales > 500
```

```
Spark.sql("SELECT * FROM orders WHERE Sales > 500").show() ##SQL equivalent
```